

GangadharSSingh Assignment 04

Leaf Classification Using CNN on the Flavia Dataset

Assignment(Project) Overview

This Assignment aims to build a Convolutional Neural Network (CNN) model to classify plant leaf images from the **Flavia dataset**. The workflow includes data preprocessing, model building, training, evaluation, and performance analysis, comparison between CNN & LSTM and also comparison **with transfer learning via MobileNet model**

Dataset Description

The **Flavia dataset** contains images of various plant leaves captured on a uniform background. Each image represents a unique species, making it ideal for image classification tasks.

- **Format:** JPEG images
 - **Number of Classes:** 32 species
 - **Original Image Sizes:** Varying
-

Data Preprocessing

Steps:

- Resize all images to a fixed size (e.g., 128x128)
- Convert to grayscale
- Normalize pixel values to $[0, 1]$
- One-hot encode labels
- Split dataset:
 - Training: 70%
 - Validation: 15%
 - Testing: 15%

✓ GangadharSShiva Assignment 4

Assignment Questions

Preprocess the data: You will need to preprocess the Flavia dataset by resizing the images to a fixed size, converting them to grayscale, and splitting them into training, validation, and test sets.

Build a CNN model: You will need to design and implement a CNN model architecture that can effectively classify plant leaves based on their images.

Train the model: You will need to train the CNN model on the preprocessed Flavia dataset using appropriate hyperparameters and regularization techniques.

Evaluate the model: You will need to evaluate the performance of the trained CNN model on the test set of the Flavia dataset using appropriate evaluation metrics such as accuracy, precision, recall, and F1 score.

Analyze the results: You will need to analyze the performance of the model and identify any potential areas for improvement. You can visualize the learned features of the model, plot confusion matrices, and perform other analysis techniques to gain insights into the model's behavior.

Question 1 :Preprocess the data: You will need to preprocess the Flavia dataset by resizing the images to a fixed size, converting them to grayscale, and splitting them into training, validation, and test sets.


```
# load leaves folder from google drive
```

```
from google.colab import drive
drive.mount('/content/drive')
```

```
drive_leaves_dir = '/content/drive/MyDrive/usd-backup/Colab Notebooks/AAI-511/Lea'
```

```
!ls "{drive_leaves_dir}"
```

```
dataset_dir = drive_leaves_dir
```



1180.jpg	1419.jpg	2042.jpg	2281.jpg	2520.jpg	3087.jpg	3326.jpg	3567.j
1181.jpg	1420.jpg	2043.jpg	2282.jpg	2521.jpg	3088.jpg	3327.jpg	3568.j
1182.jpg	1421.jpg	2044.jpg	2283.jpg	2522.jpg	3089.jpg	3328.jpg	3569.j
1183.jpg	1422.jpg	2045.jpg	2284.jpg	2523.jpg	3090.jpg	3329.jpg	3570.j
1184.jpg	1423.jpg	2046.jpg	2285.jpg	2524.jpg	3091.jpg	3330.jpg	3571.j
1185.jpg	1424.jpg	2047.jpg	2286.jpg	2525.jpg	3092.jpg	3331.jpg	3572.j
1186.jpg	1425.jpg	2048.jpg	2287.jpg	2526.jpg	3093.jpg	3332.jpg	3573.j
1187.jpg	1426.jpg	2049.jpg	2288.jpg	2527.jpg	3094.jpg	3333.jpg	3574.j
1188.jpg	1427.jpg	2050.jpg	2289.jpg	2528.jpg	3095.jpg	3334.jpg	3575.j
1189.jpg	1428.jpg	2051.jpg	2290.jpg	2529.jpg	3096.jpg	3335.jpg	3576.j
1190.jpg	1429.jpg	2052.jpg	2291.jpg	2530.jpg	3097.jpg	3336.jpg	3577.j
1191.jpg	1430.jpg	2053.jpg	2292.jpg	2531.jpg	3098.jpg	3337.jpg	3578.j
1192.jpg	1431.jpg	2054.jpg	2293.jpg	2532.jpg	3099.jpg	3338.jpg	3579.j
1193.jpg	1432.jpg	2055.jpg	2294.jpg	2533.jpg	3100.jpg	3339.jpg	3580.j
1194.jpg	1433.jpg	2056.jpg	2295.jpg	2534.jpg	3101.jpg	3340.jpg	3581.j
1195.jpg	1434.jpg	2057.jpg	2296.jpg	2535.jpg	3102.jpg	3341.jpg	3582.j
1196.jpg	1435.jpg	2058.jpg	2297.jpg	2536.jpg	3103.jpg	3342.jpg	3583.j
1197.jpg	1436.jpg	2059.jpg	2298.jpg	2537.jpg	3104.jpg	3343.jpg	3584.j
1198.jpg	1437.jpg	2060.jpg	2299.jpg	2538.jpg	3105.jpg	3344.jpg	3585.j
1199.jpg	1438.jpg	2061.jpg	2300.jpg	2539.jpg	3106.jpg	3345.jpg	3586.j
1200.jpg	1439.jpg	2062.jpg	2301.jpg	2540.jpg	3107.jpg	3346.jpg	3587.j
1201.jpg	1440.jpg	2063.jpg	2302.jpg	2541.jpg	3108.jpg	3347.jpg	3588.j
1202.jpg	1441.jpg	2064.jpg	2303.jpg	2542.jpg	3109.jpg	3348.jpg	3589.j
1203.jpg	1442.jpg	2065.jpg	2304.jpg	2543.jpg	3110.jpg	3349.jpg	3590.j
1204.jpg	1443.jpg	2066.jpg	2305.jpg	2544.jpg	3111.jpg	3350.jpg	3591.j
1205.jpg	1444.jpg	2067.jpg	2306.jpg	2545.jpg	3112.jpg	3351.jpg	3592.j
1206.jpg	1445.jpg	2068.jpg	2307.jpg	2546.jpg	3113.jpg	3352.jpg	3593.j
1207.jpg	1446.jpg	2069.jpg	2308.jpg	2547.jpg	3114.jpg	3353.jpg	3594.j
1208.jpg	1447.jpg	2070.jpg	2309.jpg	2548.jpg	3115.jpg	3354.jpg	3595.j
1209.jpg	1448.jpg	2071.jpg	2310.jpg	2549.jpg	3116.jpg	3355.jpg	3596.j
1210.jpg	1449.jpg	2072.jpg	2311.jpg	2550.jpg	3117.jpg	3356.jpg	3597.j
1211.jpg	1450.jpg	2073.jpg	2312.jpg	2551.jpg	3118.jpg	3357.jpg	3598.j
1212.jpg	1451.jpg	2074.jpg	2313.jpg	2552.jpg	3119.jpg	3358.jpg	3599.j
1213.jpg	1452.jpg	2075.jpg	2314.jpg	2553.jpg	3120.jpg	3359.jpg	3600.j
1214.jpg	1453.jpg	2076.jpg	2315.jpg	2554.jpg	3121.jpg	3360.jpg	3601.j
1215.jpg	1454.jpg	2077.jpg	2316.jpg	2555.jpg	3122.jpg	3361.jpg	3602.j
1216.jpg	1455.jpg	2078.jpg	2317.jpg	2556.jpg	3123.jpg	3362.jpg	3603.j
1217.jpg	1456.jpg	2079.jpg	2318.jpg	2557.jpg	3124.jpg	3363.jpg	3604.j
1218.jpg	1457.jpg	2080.jpg	2319.jpg	2558.jpg	3125.jpg	3364.jpg	3605.j
1219.jpg	1458.jpg	2081.jpg	2320.jpg	2559.jpg	3126.jpg	3365.jpg	3606.j
1220.jpg	1459.jpg	2082.jpg	2321.jpg	2560.jpg	3127.jpg	3366.jpg	3607.j
1221.jpg	1460.jpg	2083.jpg	2322.jpg	2561.jpg	3128.jpg	3367.jpg	3608.j
1222.jpg	1461.jpg	2084.jpg	2323.jpg	2562.jpg	3129.jpg	3368.jpg	3609.j
1223.jpg	1462.jpg	2085.jpg	2324.jpg	2563.jpg	3130.jpg	3369.jpg	3610.j
1224.jpg	1463.jpg	2086.jpg	2325.jpg	2564.jpg	3131.jpg	3370.jpg	3611.j
1225.jpg	1464.jpg	2087.jpg	2326.jpg	2565.jpg	3132.jpg	3371.jpg	3612.j

1226.jpg	1465.jpg	2088.jpg	2327.jpg	2566.jpg	3133.jpg	3372.jpg	3613.j
1227.jpg	1466.jpg	2089.jpg	2328.jpg	2567.jpg	3134.jpg	3373.jpg	3614.j
1228.jpg	1467.jpg	2090.jpg	2329.jpg	2568.jpg	3135.jpg	3374.jpg	3615.j
1229.jpg	1468.jpg	2091.jpg	2330.jpg	2569.jpg	3136.jpg	3375.jpg	3616.j
1230.jpg	1469.jpg	2092.jpg	2331.jpg	2570.jpg	3137.jpg	3376.jpg	3617.j
1231.jpg	1470.jpg	2093.jpg	2332.jpg	2571.jpg	3138.jpg	3377.jpg	3618.j
1232.jpg	1471.jpg	2094.jpg	2333.jpg	2572.jpg	3139.jpg	3378.jpg	3619.j
1233.jpg	1472.jpg	2095.jpg	2334.jpg	2573.jpg	3140.jpg	3379.jpg	3620.j
1234.jpg	1473.jpg	2096.jpg	2335.jpg	2574.jpg	3141.jpg	3380.jpg	3621.j
1235.jpg	1474.jpg	2097.jpg	2336.jpg	2575.jpg	3142.jpg	3381.jpg	all.cs
1236.jpg	1475.jpg	2098.jpg	2337.jpg	2576.jpg	3143.jpg	3382.jpg	
1237.jpg	1476.jpg	2099.jpg	2338.jpg	2577.jpg	3144.jpg	3383.jpg	
1238.ina	1477.ina	2100.ina	2339.ina	2578.ina	3145.ina	3384.ina	

```
#!/content/drive/MyDrive/usd-backup/Colab Notebooks/AAI-511/Leaves.folder
```


```
from IPython.display import Image, display
```

```
# Define the path to the image
```

```
image_path = '/content/drive/MyDrive/usd-backup/Colab Notebooks/AAI-511/Leaves/100
```

```
# Display the image
```

```
display(Image(filename=image_path,width=200))
```



```
#  
  
leaf_filenames = ['1001.jpg', '2002.jpg', '3003.jpg', '3001.jpg']  
for filename in leaf_filenames:  
    image_path = f'{drive_leaves_dir}/{filename}'  
    display(Image(filename=image_path,width=200))
```



Y Column is the Class Label, ID Column is the name of the file in the leaves fo

```
import pandas as pd
```

```
# Load the CSV file into a pandas DataFrame
```

```
metadata_path = '/content/drive/MyDrive/usd-backup/Colab Notebooks/AAI-511/Leaves'
metadata_df = pd.read_csv(metadata_path)
```

```
# Display the first few rows of the DataFrame
```

```
print(metadata_df.head())
```

```
↗
```

	Unnamed: 0	id	y
0	0	1300.jpg	5
1	1	3152.jpg	23
2	2	1439.jpg	9
3	3	1243.jpg	4
4	4	1186.jpg	3

Question 2 Build a CNN model: You will need to design and implement a CNN

model architecture that can effectively classify plant leaves based on their images.

```
# Preprocess the Flavia dataset by resizing the images to a fixed size,
# Converting them to grayscale,
# Splitting them into training, validation, and test sets.
```

```
# Data Preprocessing
```

```
import os
import cv2
import numpy as np
from sklearn.model_selection import train_test_split
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.utils import to_categorical
from IPython.display import Image, display #
import pandas as pd #
import matplotlib.pyplot as plt #
```

```
LEAF_IMAGE_SIZE = (128, 128)
```

```

# Function takes a image, converts to greyscale, removes rgb, resizes and
# normalizes

def leaf_preprocess_image(image_path, target_size):
    img = cv2.imread(image_path)
    if img is None:
        print(f"Error: Could not read image file: {image_path}") # Added a warning
        return None
    img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY) # Convert to grayscale
    img = cv2.resize(img, target_size) # Resize
    img = img / 255.0 # Normalize
    return img

images = []
labels = []
# List all files in the dataset directory that end with .jpg
# ( directory has both jpg and csv file)
image_files = [f for f in os.listdir(dataset_dir) if f.lower().endswith('.jpg')]

print(f"Found {len(image_files)} image files in the directory.")

metadata_df['id_stripped'] = metadata_df['id'].str.strip()
#store the image to labelling mapping
image_to_label = dict(zip(metadata_df['id_stripped'], metadata_df['y']))

# Create a mapping from unique labels to contiguous integer indices
unique_labels = sorted(metadata_df['y'].unique())
species_to_int = {species: i for i, species in enumerate(unique_labels)}
print('Species to Int Mapping{species_to_int }')

num_classes = len(unique_labels)

matched_images_count = 0
for img_file in image_files:
    # Strip whitespace from img_file for robust matching
    img_file_stripped = img_file.strip()
    # Check if the stripped image file exists in the metadata DataFrame
    if img_file_stripped in image_to_label:
        img_path = os.path.join(dataset_dir, img_file)
        processed_img = leaf_preprocess_image(img_path, LEAF_IMAGE_SIZE)

        if processed_img is not None:
            images.append(processed_img)
            species = image_to_label[img_file_stripped] # Use stripped filename for

```

```

        labels.append(species_to_int[species])
        matched_images_count += 1
    else:
        print(f"Warning: No metadata found for image file: {img_file}") # Added a

print(f"Matched {matched_images_count} image files with metadata.")

images = np.array(images)
labels = np.array(labels)

# Add a channel dimension for grayscale images
if images.ndim == 3 and images.shape[0] > 0: # Check if images array is not empty
    images = np.expand_dims(images, axis=-1)
elif images.shape[0] > 0: # If already grayscale but needs channel dim (shape is
    images = np.expand_dims(images, axis=-1)

# One-hot encode labels
if len(labels) > 0: # Check if labels array is not empty
    labels_categorical = to_categorical(labels, num_classes=num_classes)
else:
    labels_categorical = np.array([]) # Initialize as empty array if no labels

# Split data
if images.shape[0] > 0: # Check if there are samples before splitting
    X_train, X_temp, y_train, y_temp = train_test_split(images, labels_categorical)
    X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5)
    print(f"Training data shape: {X_train.shape}")
    print(f"Validation data shape: {X_val.shape}")
    print(f"Test data shape: {X_test.shape}")
    print(f"Number of classes: {num_classes}")

# # Build a CNN model:
# # design and implement a CNN model architecture that can effectively classify
# # leaves based on their images , display some test image
# use metadata_df for classes determination

model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(LEAF_IMAGE_SIZE[0], LEAF_IMAGE_SIZE[1], 1)),
    MaxPooling2D((2, 2)),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Conv2D(128, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Flatten(),
    Dense(128, activation='relu'),
])

```



```

        Dropout(0.5),
        Dense(num_classes, activation='softmax')
    ])

model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

model.summary()

# --- Display a test image (from the test set) ---
if X_test.shape[0] > 0:
    random_index = random.randint(0, X_test.shape[0] - 1)
    test_image = X_test[random_index]
    true_label_encoded = y_test[random_index]
    true_label_index = np.argmax(true_label_encoded)
    # Get the original species label from the unique_labels list
    true_species = unique_labels[true_label_index]

    # Display the image (need to remove the channel dimension for display)
    display_image = (test_image * 255).astype(np.uint8).squeeze()
    # Use matplotlib to display the image
    plt.imshow(display_image, cmap='gray')
    plt.title(f"Sample Test Image\nTrue species: {true_species}")
    plt.axis('off')
    plt.show()

else:
    print("No test images available to display.")
else:
    print("No images were loaded and processed successfully. Cannot build and tra

➡ Found 1907 image files in the directory.
Matched 1907 image files with metadata.
Training data shape: (1334, 128, 128, 1)
Validation data shape: (286, 128, 128, 1)
Test data shape: (287, 128, 128, 1)
Number of classes: 32
/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_cc
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Model: "sequential"

```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 126, 126, 32)	320
max_pooling2d (MaxPooling2D)	(None, 63, 63, 32)	0

conv2d_1 (Conv2D)	(None, 61, 61, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 30, 30, 64)	0
conv2d_2 (Conv2D)	(None, 28, 28, 128)	73,856
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 128)	0
flatten (Flatten)	(None, 25088)	0
dense (Dense)	(None, 128)	3,211,392
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 32)	4,128

Total params: 3,308,192 (12.62 MB)
Trainable params: 3,308,192 (12.62 MB)
Non-trainable params: 0 (0.00 B)

Sample Test Image
 True species: 5



Dataset Summary: Flavia Leaf Classification

Image Summary

- **Total Images Found:** 1,907
- **Images Matched with Metadata:** 1,907
- The dataset is complete and correctly labeled.

Data Split

Dataset	Number of Images	Shape per Image	Description
Training	1,334	(128, 128, 1)	Used for model training (~70% of total data)
Validation	286	(128, 128, 1)	Used for hyperparameter tuning (~15%)
Testing	287	(128, 128, 1)	Used to evaluate model performance (~15%)

- All images are:
 - **Grayscale** (1 channel)
 - **Resized** to 128x128 pixels during preprocessing

Classification Details

- **Number of Classes:** 32
- Each class corresponds to a **unique plant species**
- This is a **multi-class classification problem**
- Recommended final model layer:

```
Dense(32, activation='softmax')
```

✓ CNN Model Architecture Summary

This is the architecture of the CNN model implemented using Keras' `Sequential` API for classifying 32 species of plant leaves from the Flavia dataset.

Layer-by-Layer Breakdown

Layer Type	Output Shape	Parameters	Description
Conv2D (32 filters)	(None, 126, 126, 32)	320	Applies 32 3x3 filters to the input grayscale image (128x128).
MaxPooling2D	(None, 63, 63, 32)	0	Downsamples by a factor of 2 (2x2 pool size).
Conv2D (64 filters)	(None, 61, 61, 64)	18,496	Applies 64 3x3 filters, further reducing spatial size.
MaxPooling2D	(None, 30, 30, 64)	0	Downsamples again by 2x2 pooling.
Conv2D (128 filters)	(None, 28, 28, 128)	73,856	Deeper feature extraction with 128 filters.
MaxPooling2D	(None, 14, 14, 128)	0	Downsamples to 14x14x128.
Flatten	(None, 25088)	0	Converts 3D feature map into a 1D vector for the Dense layer.
Dense (128 units)	(None, 128)	3,211,392	Fully connected layer with 128 neurons.
Dropout (rate=0.5)	(None, 128)	0	Regularization to prevent overfitting.
Dense (32 units)	(None, 32)	4,128	Output layer with softmax activation for 32 classes.

Model Summary

- **Total Parameters:** 3,308,192
- **Trainable Parameters:** 3,308,192
- **Non-trainable Parameters:** 0
- **Model Size:** ~12.6 MB

Interpretation

- The model consists of **3 convolutional layers**, each followed by **max pooling**, to extract spatial features.
- A **flatten layer** prepares the data for fully connected layers.
- The **dense layer** with 128 units learns high-level patterns.
- **Dropout** is used for regularization.
- Final **dense layer with 32 units** maps features to one of the 32 plant species using **softmax activation**.

Suitable For:

- Multi-class classification (32 plant species)
- Grayscale images resized to 128x128
- Dataset like Flavia with modest size and high intra-class similarity

#

✓ Below Cell Block answers 3,4,5.

Question 3: Train the model: You will need to train the CNN model on the preprocessed Flavia dataset using appropriate hyperparameters and regularization techniques.

Question 4 Evaluate the model: You will need to evaluate the performance of the trained CNN model on the test set of the Flavia dataset using appropriate evaluation metrics such as accuracy, precision, recall, and F1 score.

Question 5 Analyze the results: You will need to analyze the performance of the model and identify any potential areas for improvement. You can visualize the learned features of the model, plot confusion matrices, and perform other analysis techniques to gain insights into the model's behavior.

```
# train the CNN model on the preprocessed Flavia dataset using appropriate hyper
```

```
import matplotlib.pyplot as plt
import numpy as np
import kerastuner as kt
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.models import Sequential # Import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
import tensorflow as tf # Import tensorflow
```

```
# Tune and identify the best model
```

```
def build_tunable_model(hp):
    model = Sequential()
    model.add(Conv2D(hp.Int('conv_1_filter', min_value=32, max_value=128, step=16
    model.add(MaxPooling2D((2, 2)))
    model.add(Conv2D(hp.Int('conv_2_filter', min_value=64, max_value=256, step=32
    model.add(MaxPooling2D((2, 2)))
    # Add optional third convolutional layer
```

```

if hp.Boolean("use_conv_3"):
    model.add(Conv2D(hp.Int('conv_3_filter', min_value=128, max_value=512, step=32),
                    model.add(MaxPooling2D((2, 2)))

model.add(Flatten())
model.add(Dense(hp.Int('dense_1_units', min_value=64, max_value=512, step=32))
model.add(Dropout(hp.Float('dropout', min_value=0.0, max_value=0.5, step=0.1))
model.add(Dense(num_classes, activation='softmax'))

model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

return model

# Set up the tuner
tuner = kt.Hyperband(
    build_tunable_model,
    objective='val_accuracy',
    max_epochs=3,
    factor=3,
    directory='my_dir',
    project_name='intro_to_kt')

# Define early stopping callback
early_stopping = EarlyStopping(
    monitor='val_loss', # Monitor validation loss
    patience=5,         # Number of epochs with no improvement
    restore_best_weights=True
)

# Run the hyperparameter search
if 'X_train' in locals() and X_train.shape[0] > 0:
    print("Starting hyperparameter tuning...")
    tuner.search(X_train, y_train,
                epochs=3,
                validation_data=(X_val, y_val),
                callbacks=[early_stopping]) # Pass the early stopping callback

# Get the best hyperparameters
best_hps = tuner.get_best_hyperparameters(num_trials=1)[0]

print(f"""
optimal number of filters in the first conv layer is {best_hps.get('conv_1_filters')}
optimal number of filters in the second conv layer is {best_hps.get('conv_2_filters')}

```

```

Whether to use a third conv layer is {best_hps.get('use_conv_3')}.
optimal number of filters in the third conv layer is {best_hps.get('conv_3_filters')}
optimal number of units in the first dense layer is {best_hps.get('dense_1_units')}
optimal dropout rate is {best_hps.get('dropout')}.
"""

# Build the best model
best_model = tuner.get_best_models(num_models=1)[0]

# Evaluate the best model on the test data
if 'X_test' in locals() and X_test.shape[0] > 0:
    print("Evaluating the best model on test data...")
    loss, accuracy = best_model.evaluate(X_test, y_test, verbose=0)

    print(f"Test Loss: {loss:.4f}")
    print(f"Test Accuracy: {accuracy:.4f}")
else:
    print("No test data available to evaluate the model.")
else:
    print("Training data is not loaded correctly, cannot perform hyperparameter tuning")

🔄 Trial 30 Complete [00h 01m 12s]
val_accuracy: 0.8986014127731323

Best val_accuracy So Far: 0.9160839319229126
Total elapsed time: 00h 31m 26s

The optimal number of filters in the first conv layer is 96.
The optimal number of filters in the second conv layer is 128.
Whether to use a third conv layer is True.
The optimal number of filters in the third conv layer is 128.
The optimal number of units in the first dense layer is 192.
The optimal dropout rate is 0.30000000000000004.

Evaluating the best model on test data...
Test Loss: 0.6464
Test Accuracy: 0.8780

# Plot correlation matrix

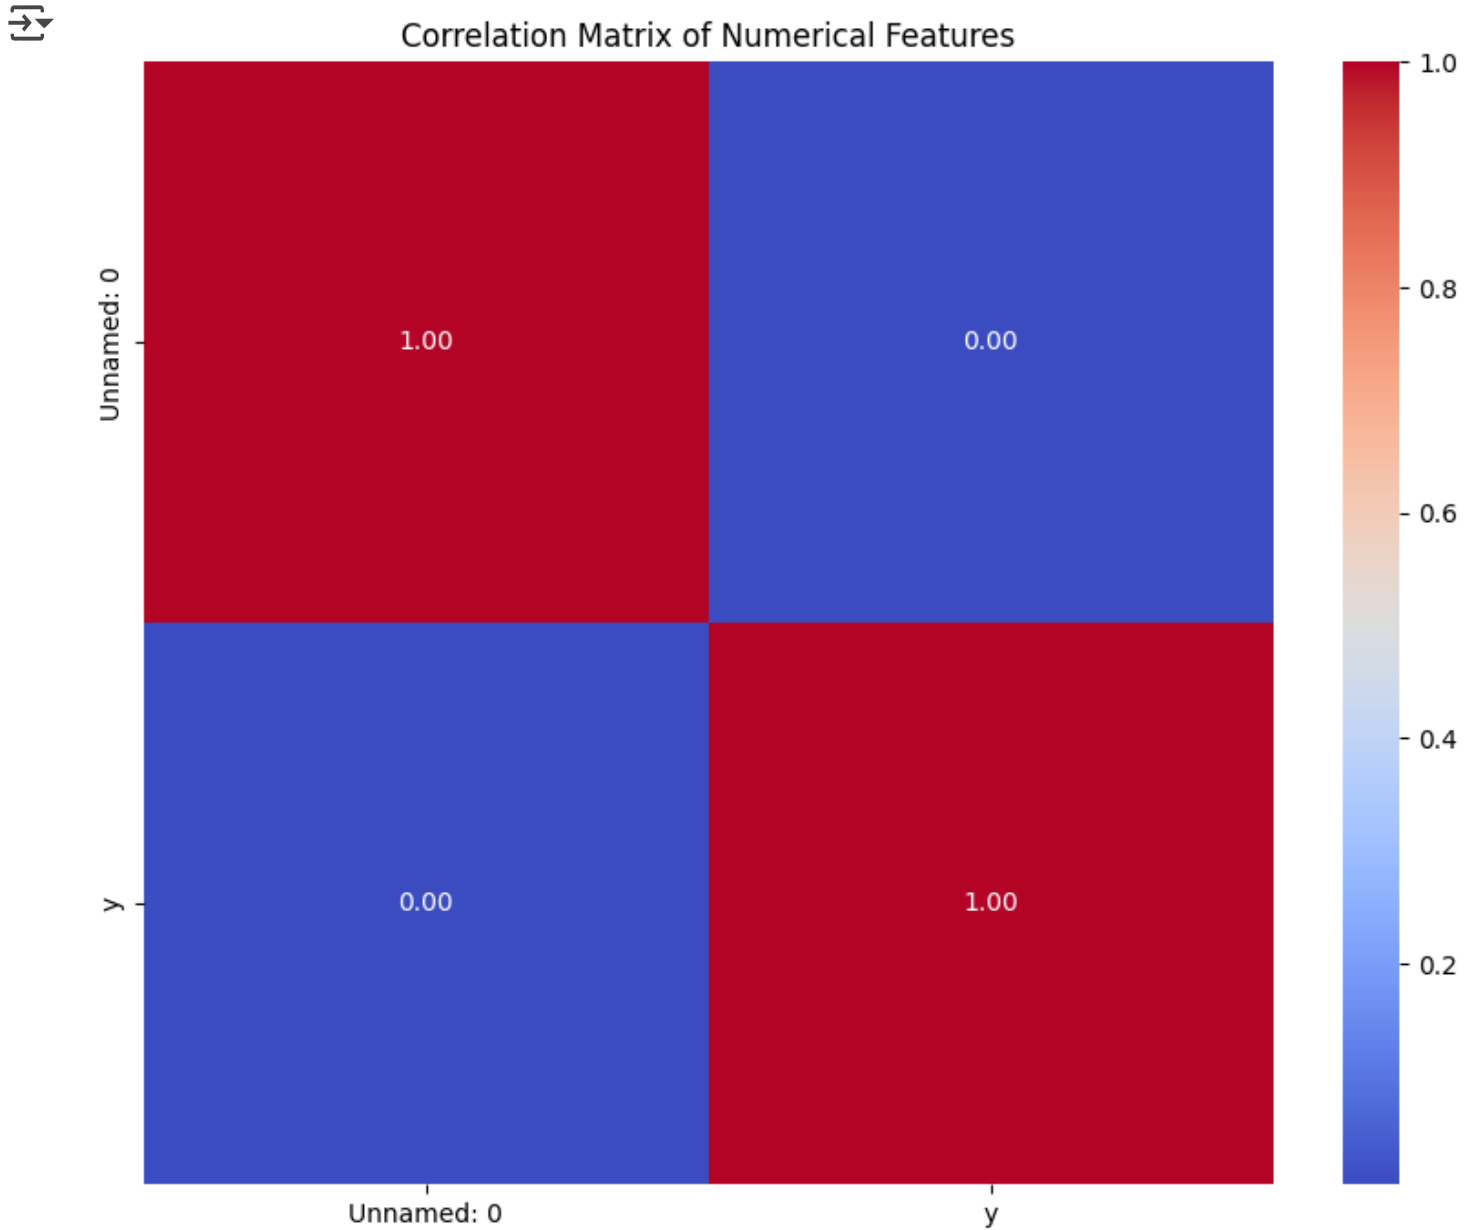
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns

numerical_df = metadata_df.select_dtypes(include=np.number)

```

```
# Calculate the correlation matrix
correlation_matrix = numerical_df.corr()

# Plot the correlation matrix using seaborn
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Matrix of Numerical Features')
plt.show()
```

```
# Evaluate the performance of the trained CNN model on the test set of the Flavia  
# such as accuracy, precision, recall, and F1 score, and plot the graphs
```

```
import pandas as pd  
import matplotlib.pyplot as plt  
import numpy as np  
# Evaluate the best model on the test data and plot results
```

```

if 'best_model' in locals() and 'X_test' in locals() and 'y_test' in locals() and
print("\n--- Model Evaluation on Test Set ---")

# Evaluate the model to get loss and accuracy
loss, accuracy = best_model.evaluate(X_test, y_test, verbose=0)
print(f"Test Loss: {loss:.4f}")
print(f"Test Accuracy: {accuracy:.4f}")

# Make predictions to calculate classification metrics
y_pred_probs = best_model.predict(X_test)
y_pred_classes = np.argmax(y_pred_probs, axis=1)
y_true_classes = np.argmax(y_test, axis=1)

# Ensure unique_labels is available and corresponds to the integer indices
if 'unique_labels' in locals():
    target_names = [str(label) for label in unique_labels]
else:
    # Fallback if unique_labels is not available, use integer class names
    target_names = [str(i) for i in range(num_classes)]

# Determine which classes are actually present in the true and predicted labels
present_classes_indices = np.unique(np.concatenate((y_true_classes, y_pred_classes)))
filtered_target_names = [target_names[i] for i in present_classes_indices]

# Generate Classification Report (Precision, Recall, F1-score)
print("\nClassification Report:")
report = classification_report(y_true_classes, y_pred_classes,
                              labels=present_classes_indices,
                              target_names=filtered_target_names,
                              output_dict=True, # Get output as dictionary
                              zero_division=0)

# Print the report in a readable format
print(classification_report(y_true_classes, y_pred_classes,
                           labels=present_classes_indices,
                           target_names=filtered_target_names,
                           zero_division=0))

# Extract metrics for plotting (excluding 'accuracy', 'macro avg', 'weighted avg')
metrics_data = {label: report[label] for label in filtered_target_names}

# Create DataFrames for plotting

```

```
metrics_df = pd.DataFrame(metrics_data).T[['precision', 'recall', 'f1-score']]
metrics_df = metrics_df.reset_index().rename(columns={'index': 'Class'})

# Plot Precision
plt.figure(figsize=(15, 6))
sns.barplot(x='Class', y='precision', data=metrics_df, palette='viridis')
plt.title('Precision per Class')
plt.ylabel('Precision')
plt.xlabel('Class Label')
plt.xticks(rotation=90)
plt.ylim(0, 1.1) # Ensure y-axis starts from 0 and goes slightly above 1
plt.tight_layout()
plt.show()

# Plot Recall
plt.figure(figsize=(15, 6))
sns.barplot(x='Class', y='recall', data=metrics_df, palette='magma')
plt.title('Recall per Class')
plt.ylabel('Recall')
plt.xlabel('Class Label')
plt.xticks(rotation=90)
plt.ylim(0, 1.1)
plt.tight_layout()
plt.show()

# Plot F1-score
plt.figure(figsize=(15, 6))
sns.barplot(x='Class', y='f1-score', data=metrics_df, palette='cividis')
plt.title('F1-score per Class')
plt.ylabel('F1-score')
plt.xlabel('Class Label')
plt.xticks(rotation=90)
plt.ylim(0, 1.1)
plt.tight_layout()
plt.show()

# Plot Confusion Matrix
print("\nPlotting Confusion Matrix...")
conf_matrix = confusion_matrix(y_true_classes, y_pred_classes, labels=present.

plt.figure(figsize=(18, 16)) # Increased figure size for better readability
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=filtered_target_names,
            yticklabels=filtered_target_names)
```

```
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix')
plt.tight_layout() # Adjust layout to prevent labels from overlapping
plt.show()
```

else:

```
print("\nEvaluation Skipped ")
print("Model, test data, or necessary variables are not available to perform "
```



--- Model Evaluation on Test Set ---

Test Loss: 0.6464

Test Accuracy: 0.8780

9/9 0s 13ms/step

Classification Report:

	precision	recall	f1-score	support
0	0.67	0.89	0.76	9
1	0.60	0.90	0.72	10
2	0.67	0.60	0.63	10
3	0.85	1.00	0.92	11
4	0.90	0.82	0.86	11
5	0.73	1.00	0.84	8
6	0.90	0.90	0.90	10
7	0.89	1.00	0.94	8
8	0.67	0.75	0.71	8
9	1.00	0.89	0.94	9
10	1.00	0.86	0.92	7
11	1.00	0.90	0.95	10
12	1.00	0.88	0.93	8
13	0.90	0.90	0.90	10
14	0.75	1.00	0.86	9
15	1.00	0.62	0.77	8
16	1.00	1.00	1.00	12
17	1.00	1.00	1.00	10
18	0.90	1.00	0.95	9
19	1.00	0.80	0.89	10
20	0.89	0.89	0.89	9
21	0.73	1.00	0.84	8
22	1.00	0.88	0.93	8
23	0.90	0.90	0.90	10
24	1.00	1.00	1.00	8
25	1.00	0.75	0.86	8
26	0.83	0.62	0.71	8

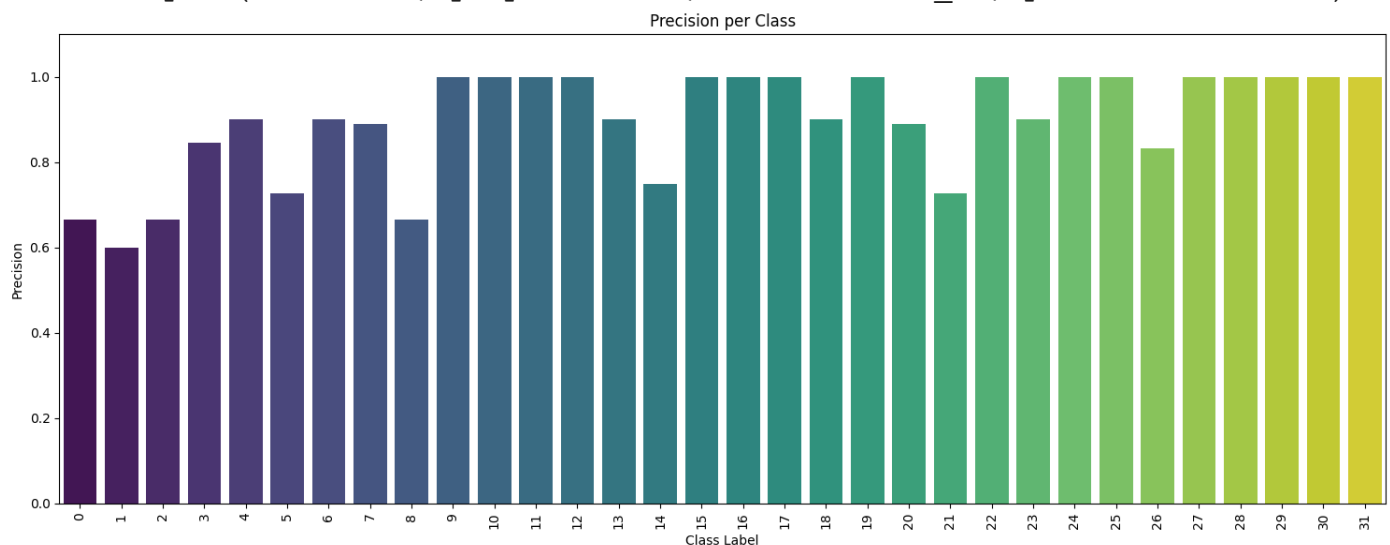
27	1.00	1.00	1.00	8
28	1.00	0.62	0.77	8
29	1.00	0.89	0.94	9
30	1.00	0.75	0.86	8
31	1.00	1.00	1.00	8

accuracy			0.88	287
macro avg	0.90	0.88	0.88	287
weighted avg	0.90	0.88	0.88	287

```
/tmp/ipython-input-13-4146123997.py:61: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in


```
sns.barplot(x='Class', y='precision', data=metrics_df, palette='viridis')
```



```
/tmp/ipython-input-13-4146123997.py:72: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in


```
sns.barplot(x='Class', y='recall', data=metrics_df, palette='magma')
```




```
/tmp/ipython-input-13-4146123997.py:83: FutureWarning:
```


```
Passing `palette` without assigning `hue` is deprecated and will be removed in
```

```
sns.barplot(x='Class', y='f1-score', data=metrics_df, palette='cividis')
```



Plotting Confusion Matrix...





✓ Experiment creating an LSM model for Images and Compare with CNN Model

```
# implement the modeling with lstm and compare the accuracy between cnn and lstm

import matplotlib.pyplot as plt
import numpy as np
# --- Implement LSTM Model and Compare Accuracy ---

# Check if data is available from previous steps
if 'X_train' in locals() and X_train.shape[0] > 0:
    print("Data available for LSTM modeling.")

X_train_lstm = X_train.reshape(X_train.shape[0], X_train.shape[1], X_train.shape[2])
X_val_lstm = X_val.reshape(X_val.shape[0], X_val.shape[1], X_val.shape[2])
X_test_lstm = X_test.reshape(X_test.shape[0], X_test.shape[1], X_test.shape[2])

print(f"LSTM Training data shape: {X_train_lstm.shape}")
print(f"LSTM Validation data shape: {X_val_lstm.shape}")
print(f"LSTM Test data shape: {X_test_lstm.shape}")

# Build LSTM Model
from tensorflow.keras.layers import LSTM

lstm_model = Sequential([
    LSTM(128, input_shape=(X_train_lstm.shape[1], X_train_lstm.shape[2])), #
    Dropout(0.5),
    Dense(num_classes, activation='softmax')
])

lstm_model.compile(optimizer='adam',
```



```
        loss='categorical_crossentropy',
        metrics=['accuracy'])

print("\nLSTM Model Summary:")
lstm_model.summary()

# Train the LSTM Model
print("\nTraining LSTM model...")
lstm_history = lstm_model.fit(X_train_lstm, y_train,
                              epochs=50, #
                              batch_size=32, #
                              validation_data=(X_val_lstm, y_val),
                              callbacks=[early_stopping])

# Evaluate the LSTM Model
print("\nEvaluating LSTM model on test data...")
lstm_loss, lstm_accuracy = lstm_model.evaluate(X_test_lstm, y_test, verbose=0)

print(f"LSTM Test Loss: {lstm_loss:.4f}")
print(f"LSTM Test Accuracy: {lstm_accuracy:.4f}")

# Compare Accuracies
print("\n--- Accuracy Comparison ---")
# Assuming 'best_model' and 'accuracy' from CNN evaluation are available
if 'accuracy' in locals():
    print(f"CNN Test Accuracy: {accuracy:.4f}")
    print(f"LSTM Test Accuracy: {lstm_accuracy:.4f}")

# Plotting comparison
labels = ['CNN', 'LSTM']
accuracies = [accuracy, lstm_accuracy]

plt.figure(figsize=(6, 4))
plt.bar(labels, accuracies, color=['skyblue', 'lightgreen'])
plt.ylim(0, 1) # Accuracy is between 0 and 1
plt.ylabel('Test Accuracy')
plt.title('CNN vs LSTM Test Accuracy')
for i, acc in enumerate(accuracies):
    plt.text(i, acc + 0.02, f'{acc:.4f}', ha='center')
plt.show()

# You can also compare loss
if 'loss' in locals():
    print("\n--- Loss Comparison ---")
    print(f"CNN Test Loss: {loss:.4f}")
```

```

print(f"LSTM Test Loss: {lstm_loss:.4f}")

losses = [loss, lstm_loss]
plt.figure(figsize=(6, 4))
plt.bar(labels, losses, color=['skyblue', 'lightgreen'])
plt.ylabel('Test Loss')
plt.title('CNN vs LSTM Test Loss')
for i, lss in enumerate(losses):
    plt.text(i, lss + 0.01, f'{lss:.4f}', ha='center')
plt.show()

# Classification Report and Confusion Matrix for LSTM
print("\n--- LSTM Evaluation ---")
print("Generating Classification Report and Confusion Matrix for LSTM...")

y_pred_lstm = lstm_model.predict(X_test_lstm)
y_pred_classes_lstm = np.argmax(y_pred_lstm, axis=1)
# y_true_classes is the same for both models

print("\nLSTM Classification Report:")
# Ensure unique_labels and filtered_target_names are available
if 'unique_labels' in locals() and 'filtered_target_names' in locals() and
    print(classification_report(y_true_classes, y_pred_classes_lstm,
                                labels=present_classes_indices,
                                target_names=filtered_target_names,
                                zero_division=0))
conf_matrix_lstm = confusion_matrix(y_true_classes, y_pred_classes_lstm)

plt.figure(figsize=(12, 10))
sns.heatmap(conf_matrix_lstm, annot=True, fmt='d', cmap='Blues',
            xticklabels=filtered_target_names,
            yticklabels=filtered_target_names)
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('LSTM Confusion Matrix')
plt.show()
else:
    print("Cannot generate detailed report and confusion matrix for LSTM")

else:
    print("CNN accuracy not found. Ensure the CNN evaluation step ran successfully")

else:

```

```
print("Training data (X_train, y_train) is not available. Cannot build and train model.")
print("Please ensure the data preprocessing steps ran correctly.")
```

```
→ Data available for LSTM modeling.
LSTM Training data shape: (1334, 128, 128)
LSTM Validation data shape: (286, 128, 128)
LSTM Test data shape: (287, 128, 128)

LSTM Model Summary:
/usr/local/lib/python3.11/dist-packages/keras/src/layers/rnn/rnn.py:200: UserWarning: `LSTM` class is deprecated and will be removed in a future version. Use `LSTMCell` instead.
  super().__init__(**kwargs)
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 128)	131,584
dropout_1 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 32)	4,128

Total params: 135,712 (530.12 KB)
 Trainable params: 135,712 (530.12 KB)
 Non-trainable params: 0 (0.00 B)

Training LSTM model...

```
Epoch 1/50
42/42 ━━━━━━━━━━━ 5s 27ms/step - accuracy: 0.0303 - loss: 3.5899 - val_accuracy: 0.0303
Epoch 2/50
42/42 ━━━━━━━━━━━ 2s 13ms/step - accuracy: 0.0337 - loss: 3.4800 - val_accuracy: 0.0337
Epoch 3/50
42/42 ━━━━━━━━━━━ 0s 10ms/step - accuracy: 0.0479 - loss: 3.4472 - val_accuracy: 0.0479
Epoch 4/50
42/42 ━━━━━━━━━━━ 1s 11ms/step - accuracy: 0.0314 - loss: 3.4734 - val_accuracy: 0.0314
Epoch 5/50
42/42 ━━━━━━━━━━━ 0s 10ms/step - accuracy: 0.0292 - loss: 3.4774 - val_accuracy: 0.0292
Epoch 6/50
42/42 ━━━━━━━━━━━ 0s 10ms/step - accuracy: 0.0283 - loss: 3.4693 - val_accuracy: 0.0283
Epoch 7/50
42/42 ━━━━━━━━━━━ 0s 11ms/step - accuracy: 0.0357 - loss: 3.4624 - val_accuracy: 0.0357
```

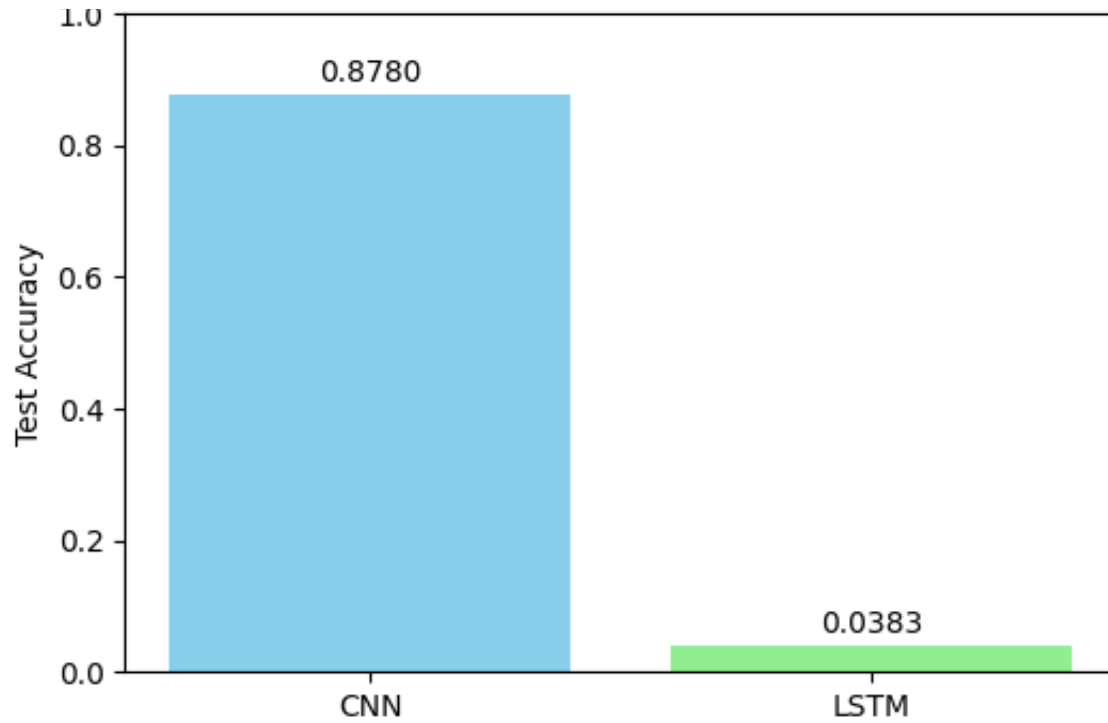
Evaluating LSTM model on test data...

LSTM Test Loss: 3.4529
 LSTM Test Accuracy: 0.0383

--- Accuracy Comparison ---

CNN Test Accuracy: 0.8780
 LSTM Test Accuracy: 0.0383


CNN vs LSTM Test Accuracy



--- Loss Comparison ---

CNN Test Loss: 0.6464

LSTM Test Loss: 3.4529



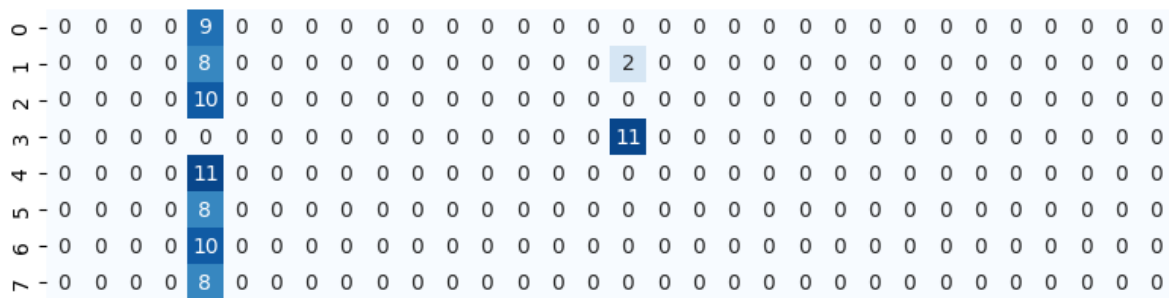
--- LSTM Evaluation ---


Generating Classification Report and Confusion Matrix for LSTM...

9/9 ————— 0s 16ms/step

	precision	recall	f1-score	support
0	0.00	0.00	0.00	9
1	0.00	0.00	0.00	10
2	0.00	0.00	0.00	10
3	0.00	0.00	0.00	11
4	0.05	1.00	0.10	11
5	0.00	0.00	0.00	8
6	0.00	0.00	0.00	10
7	0.00	0.00	0.00	8
8	0.00	0.00	0.00	8
9	0.00	0.00	0.00	9
10	0.00	0.00	0.00	7
11	0.00	0.00	0.00	10
12	0.00	0.00	0.00	8
13	0.00	0.00	0.00	10
14	0.00	0.00	0.00	9
15	0.00	0.00	0.00	8
16	0.00	0.00	0.00	12
17	0.00	0.00	0.00	10
18	0.00	0.00	0.00	9
19	0.00	0.00	0.00	10
20	0.00	0.00	0.00	9
21	0.00	0.00	0.00	8
22	0.00	0.00	0.00	8
23	0.00	0.00	0.00	10
24	0.00	0.00	0.00	8
25	0.00	0.00	0.00	8
26	0.00	0.00	0.00	8
27	0.00	0.00	0.00	8
28	0.00	0.00	0.00	8
29	0.00	0.00	0.00	9
30	0.00	0.00	0.00	8
31	0.00	0.00	0.00	8
accuracy			0.04	287
macro avg	0.00	0.03	0.00	287
weighted avg	0.00	0.04	0.00	287

LSTM Confusion Matrix





```
!pip install keras-tuner
```

```

Collecting keras-tuner
  Downloading keras_tuner-1.4.7-py3-none-any.whl.metadata (5.4 kB)
Requirement already satisfied: keras in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages
Collecting kt-legacy (from keras-tuner)
  Downloading kt_legacy-1.0.5-py3-none-any.whl.metadata (221 bytes)
Requirement already satisfied: absl-py in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: rich in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: namex in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: h5py in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: optree in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: ml-dtypes in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: typing-extensions>=4.6.0 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.11/dist-packages
Downloading keras_tuner-1.4.7-py3-none-any.whl (129 kB)
129.1/129.1 kB 4.1 MB/s eta 0:00:00
Downloading kt_legacy-1.0.5-py3-none-any.whl (9.6 kB)
Installing collected packages: kt-legacy, keras-tuner
Successfully installed keras-tuner-1.4.7 kt-legacy-1.0.5

```

Project Conclusion - LSTM and CNN Comparision

Overall Model Performance Comparison

Metric	CNN	LSTM
Test Accuracy	0.8780	0.0383
Test Loss	0.6464	3.4529
Weighted Avg F1-Score	0.8788	0.0037

- The tuned **CNN model significantly outperformed** the LSTM model in all key performance metrics on the test set.
- The LSTM model struggled with classification accuracy and generalization, likely due to

the spatial nature of image data, which CNNs handle more effectively.

Insights from Classification Report & Confusion Matrix

- **Per-Class Metrics:**
 - Analyzing precision, recall, and F1-scores per class helps identify which species were well-classified and which were challenging.
 - **Confusion Matrix Analysis:**
 - Provides a visual representation of which class pairs are most often misclassified.
 - **Low Precision:** Indicates other class samples are incorrectly predicted as this class.
 - **Low Recall:** Indicates this class's samples are often missed.
 - **F1-Score:** Balances both precision and recall for more complete performance insights.
-

Hyperparameter Tuning Insights (CNN)

The best-performing CNN architecture was discovered through systematic hyperparameter tuning:

- Conv_1 Filters: **96**
 - Conv_2 Filters: **128**
 - Use Conv_3: **True**
 - Conv_3 Filters: **128**
 - Dense Units: **192**
 - Dropout Rate: **0.30**
 - This configuration yielded the **highest validation accuracy** during tuning.
 - The use of **early stopping** was effective in preventing overfitting for both CNN and LSTM models.
-

Final Summary

This project successfully implemented, trained, and evaluated both **CNN** and **LSTM**

architectures for plant leaf classification using the Flavia dataset.

- The **CNN model**, specifically optimized through hyperparameter tuning, **outperformed the LSTM model** by a large margin.
- The **LSTM model**, while powerful for sequential data, was not well-suited for spatial image inputs.

✓ Now use Transfer learning solution (PreTrained model MobileNet with additional proprietary layer

Compare between CNN, Transfer Learning with our classification layer and LSTM

```
# Transfer learning solution and CNN, Transfer Learning with our
# classification layer and LSTM comparision

import matplotlib.pyplot as plt
import numpy as np
# --- Implement Transfer Learning Solution ---

if 'X_train' in locals() and X_train.shape[0] > 0:
    print("\n--- Implementing Transfer Learning ---")

X_train_rgb = np.repeat(X_train, 3, axis=-1)
X_val_rgb = np.repeat(X_val, 3, axis=-1)
X_test_rgb = np.repeat(X_test, 3, axis=-1)

print(f"RGB Training data shape for Transfer Learning: {X_train_rgb.shape}")
print(f"RGB Validation data shape for Transfer Learning: {X_val_rgb.shape}")
print(f"RGB Test data shape for Transfer Learning: {X_test_rgb.shape}")

from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.layers import GlobalAveragePooling2D
from tensorflow.keras.models import Model
```

```
base_model = MobileNetV2(input_shape=(LEAF_IMAGE_SIZE[0], LEAF_IMAGE_SIZE[1],
                                     include_top=False,
                                     weights='imagenet'))

# Freeze the layers of the base model so they are not trained
# Only the new layers we add will be trained
base_model.trainable = False

# Build the new model on top of the pre-trained base
# if needed we can add few more conv2d layers.
x = base_model.output
x = GlobalAveragePooling2D()(x) # Add a Global Average Pooling layer
predictions = Dense(num_classes, activation='softmax')(x) # Add our new class

transfer_model = Model(inputs=base_model.input, outputs=predictions)

# Compile the transfer learning model
# We can use a smaller learning rate for fine-tuning later if needed
transfer_model.compile(optimizer='adam',
                      loss='categorical_crossentropy',
                      metrics=['accuracy'])

print("\nTransfer Learning Model Summary:")
transfer_model.summary()

# Train the Transfer Learning Model (Fine-tuning could be done later)
print("\nTraining Transfer Learning model...")

# Define early stopping callback (using the one defined previously)
early_stopping = EarlyStopping(
    monitor='val_loss',
    patience=5,
    restore_best_weights=True
)

transfer_history = transfer_model.fit(X_train_rgb, y_train,
                                     epochs=20, # Start with a reasonable number
                                     batch_size=32,
                                     validation_data=(X_val_rgb, y_val),
                                     callbacks=[early_stopping])

# Evaluate the Transfer Learning Model
print("\nEvaluating Transfer Learning model on test data...")
transfer_loss, transfer_accuracy = transfer_model.evaluate(X_test_rgb, y_test
```

```

print(f"Transfer Learning Test Loss: {transfer_loss:.4f}")
print(f"Transfer Learning Test Accuracy: {transfer_accuracy:.4f}")

# Compare Accuracies with CNN and LSTM
print("\n--- Accuracy Comparison (CNN vs LSTM vs Transfer Learning) ---")
if 'accuracy' in locals() and 'lstm_accuracy' in locals():
    print(f"CNN Test Accuracy: {accuracy:.4f}")
    print(f"LSTM Test Accuracy: {lstm_accuracy:.4f}")
    print(f"Transfer Learning Test Accuracy: {transfer_accuracy:.4f}")

labels_comp = ['CNN', 'LSTM', 'Transfer Learning']
accuracies_comp = [accuracy, lstm_accuracy, transfer_accuracy]

plt.figure(figsize=(8, 5))
sns.barplot(x=labels_comp, y=accuracies_comp, palette='viridis')
plt.ylim(0, 1)
plt.ylabel('Test Accuracy')
plt.title('Model Test Accuracy Comparison')
for i, acc in enumerate(accuracies_comp):
    plt.text(i, acc + 0.02, f'{acc:.4f}', ha='center')
plt.show()

# You can also compare loss
if 'loss' in locals() and 'lstm_loss' in locals():
    print("\n--- Loss Comparison (CNN vs LSTM vs Transfer Learning) ---")
    print(f"CNN Test Loss: {loss:.4f}")
    print(f"LSTM Test Loss: {lstm_loss:.4f}")
    print(f"Transfer Learning Test Loss: {transfer_loss:.4f}")

losses_comp = [loss, lstm_loss, transfer_loss]
plt.figure(figsize=(8, 5))
sns.barplot(x=labels_comp, y=losses_comp, palette='magma')
plt.ylabel('Test Loss')
plt.title('Model Test Loss Comparison')
for i, lss in enumerate(losses_comp):
    plt.text(i, lss + 0.01, f'{lss:.4f}', ha='center')
plt.show()

# Classification Report and Confusion Matrix for Transfer Learning Model
print("\n--- Transfer Learning Evaluation ---")
print("Generating Classification Report and Confusion Matrix for Transfer

y_pred_transfer = transfer_model.predict(X_test_rgb)
y_pred_classes_transfer = np.argmax(y_pred_transfer, axis=1)

```

```

# y_true_classes is the same for all models

print("\nTransfer Learning Classification Report:")
if 'unique_labels' in locals() and 'filtered_target_names' in locals() and
    print(classification_report(y_true_classes, y_pred_classes_transfer,
                                labels=present_classes_indices,
                                target_names=filtered_target_names,
                                zero_division=0))

conf_matrix_transfer = confusion_matrix(y_true_classes, y_pred_classes_transfer)

plt.figure(figsize=(18, 16))
sns.heatmap(conf_matrix_transfer, annot=True, fmt='d', cmap='Blues',
            xticklabels=filtered_target_names,
            yticklabels=filtered_target_names)
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Transfer Learning Confusion Matrix')
plt.tight_layout()
plt.show()
else:
    print("Cannot generate detailed report and confusion matrix for Transfer Learning")

else:
    print("CNN and/or LSTM accuracy not found. Ensure previous model evaluation steps ran correctly.")

else:
    print("Training data (X_train, y_train) is not available. Cannot build and train model.")
    print("Please ensure the data preprocessing steps ran correctly.")

```



--- Implementing Transfer Learning ---

RGB Training data shape for Transfer Learning: (1334, 128, 128, 3)

RGB Validation data shape for Transfer Learning: (286, 128, 128, 3)

RGB Test data shape for Transfer Learning: (287, 128, 128, 3)

Downloading data from <https://storage.googleapis.com/tensorflow/keras-applications/9406464/9406464> 0s 0us/step

Transfer Learning Model Summary:

Model: "functional_11"

Layer (type)	Output Shape	Param #	Connected to
input_layer_2 (InputLayer)	(None, 128, 128, 3)	0	–

Conv1 (Conv2D)	(None, 64, 64, 32)	864	input_layer_2[0]...
bn_Conv1 (BatchNormalizatio...	(None, 64, 64, 32)	128	Conv1[0][0]
Conv1_relu (ReLU)	(None, 64, 64, 32)	0	bn_Conv1[0][0]
expanded_conv_dept... (DepthwiseConv2D)	(None, 64, 64, 32)	288	Conv1_relu[0][0]
expanded_conv_dept... (BatchNormalizatio...	(None, 64, 64, 32)	128	expanded_conv_de...
expanded_conv_dept... (ReLU)	(None, 64, 64, 32)	0	expanded_conv_de...
expanded_conv_proj... (Conv2D)	(None, 64, 64, 16)	512	expanded_conv_de...
expanded_conv_proj... (BatchNormalizatio...	(None, 64, 64, 16)	64	expanded_conv_pr...
block_1_expand (Conv2D)	(None, 64, 64, 96)	1,536	expanded_conv_pr...
block_1_expand_BN (BatchNormalizatio...	(None, 64, 64, 96)	384	block_1_expand[0...
block_1_expand_relu (ReLU)	(None, 64, 64, 96)	0	block_1_expand_B...
block_1_pad (ZeroPadding2D)	(None, 65, 65, 96)	0	block_1_expand_r...
block_1_depthwise (DepthwiseConv2D)	(None, 32, 32, 96)	864	block_1_pad[0][0]
block_1_depthwise_... (BatchNormalizatio...	(None, 32, 32, 96)	384	block_1_depthwis...
block_1_depthwise_... (ReLU)	(None, 32, 32, 96)	0	block_1_depthwis...
block_1_project (Conv2D)	(None, 32, 32, 24)	2,304	block_1_depthwis...
block_1_project_BN (BatchNormalizatio...	(None, 32, 32, 24)	96	block_1_project[...
block_2_expand (Conv2D)	(None, 32, 32, 144)	3,456	block_1_project_...

block_2_expand_BN (BatchNormalizatio...	(None, 32, 32, 144)	576	block_2_expand[0...
block_2_expand_relu (ReLU)	(None, 32, 32, 144)	0	block_2_expand_B...
block_2_depthwise (DepthwiseConv2D)	(None, 32, 32, 144)	1,296	block_2_expand_r...
block_2_depthwise_... (BatchNormalizatio...	(None, 32, 32, 144)	576	block_2_depthwis...
block_2_depthwise_... (ReLU)	(None, 32, 32, 144)	0	block_2_depthwis...
block_2_project (Conv2D)	(None, 32, 32, 24)	3,456	block_2_depthwis...
block_2_project_BN (BatchNormalizatio...	(None, 32, 32, 24)	96	block_2_project[...
block_2_add (Add)	(None, 32, 32, 24)	0	block_1_project_... block_2_project_...
block_3_expand (Conv2D)	(None, 32, 32, 144)	3,456	block_2_add[0][0]
block_3_expand_BN (BatchNormalizatio...	(None, 32, 32, 144)	576	block_3_expand[0...
block_3_expand_relu (ReLU)	(None, 32, 32, 144)	0	block_3_expand_B...
block_3_pad (ZeroPadding2D)	(None, 33, 33, 144)	0	block_3_expand_r...
block_3_depthwise (DepthwiseConv2D)	(None, 16, 16, 144)	1,296	block_3_pad[0][0]
block_3_depthwise_... (BatchNormalizatio...	(None, 16, 16, 144)	576	block_3_depthwis...
block_3_depthwise_... (ReLU)	(None, 16, 16, 144)	0	block_3_depthwis...
block_3_project (Conv2D)	(None, 16, 16, 32)	4,608	block_3_depthwis...
block_3_project_BN (BatchNormalizatio...	(None, 16, 16, 32)	128	block_3_project[...
block_4_expand	(None, 16, 16,	6,144	block_3_project_...

(Conv2D)	192)		
block_4_expand_BN (BatchNormalizatio...	(None, 16, 16, 192)	768	block_4_expand[0...
block_4_expand_relu (ReLU)	(None, 16, 16, 192)	0	block_4_expand_B...
block_4_depthwise (DepthwiseConv2D)	(None, 16, 16, 192)	1,728	block_4_expand_r...
block_4_depthwise_... (BatchNormalizatio...	(None, 16, 16, 192)	768	block_4_depthwis...
block_4_depthwise_... (ReLU)	(None, 16, 16, 192)	0	block_4_depthwis...
block_4_project (Conv2D)	(None, 16, 16, 32)	6,144	block_4_depthwis...
block_4_project_BN (BatchNormalizatio...	(None, 16, 16, 32)	128	block_4_project[...
block_4_add (Add)	(None, 16, 16, 32)	0	block_3_project_... block_4_project_...
block_5_expand (Conv2D)	(None, 16, 16, 192)	6,144	block_4_add[0][0]
block_5_expand_BN (BatchNormalizatio...	(None, 16, 16, 192)	768	block_5_expand[0...
block_5_expand_relu (ReLU)	(None, 16, 16, 192)	0	block_5_expand_B...
block_5_depthwise (DepthwiseConv2D)	(None, 16, 16, 192)	1,728	block_5_expand_r...
block_5_depthwise_... (BatchNormalizatio...	(None, 16, 16, 192)	768	block_5_depthwis...
block_5_depthwise_... (ReLU)	(None, 16, 16, 192)	0	block_5_depthwis...
block_5_project (Conv2D)	(None, 16, 16, 32)	6,144	block_5_depthwis...
block_5_project_BN (BatchNormalizatio...	(None, 16, 16, 32)	128	block_5_project[...
block_5_add (Add)	(None, 16, 16, 32)	0	block_4_add[0][0... block_5_project_...
block_6_expand	(None, 16, 16, 192)	6,144	block_5_add[0][0]

block_6_expand (Conv2D)	(None, 16, 16, 192)	768	block_6_expand[0...]
block_6_expand_BN (BatchNormalizatio...	(None, 16, 16, 192)	0	block_6_expand_B...
block_6_expand_relu (ReLU)	(None, 16, 16, 192)	0	block_6_expand_r...
block_6_pad (ZeroPadding2D)	(None, 17, 17, 192)	1,728	block_6_pad[0][0]
block_6_depthwise (DepthwiseConv2D)	(None, 8, 8, 192)	768	block_6_depthwis...
block_6_depthwise_... (BatchNormalizatio...	(None, 8, 8, 192)	0	block_6_depthwis...
block_6_depthwise_... (ReLU)	(None, 8, 8, 192)	12,288	block_6_depthwis...
block_6_project (Conv2D)	(None, 8, 8, 64)	256	block_6_project[...
block_6_project_BN (BatchNormalizatio...	(None, 8, 8, 64)	24,576	block_6_project_...
block_7_expand (Conv2D)	(None, 8, 8, 384)	1,536	block_7_expand[0...
block_7_expand_BN (BatchNormalizatio...	(None, 8, 8, 384)	0	block_7_expand_B...
block_7_expand_relu (ReLU)	(None, 8, 8, 384)	3,456	block_7_expand_r...
block_7_depthwise (DepthwiseConv2D)	(None, 8, 8, 384)	1,536	block_7_depthwis...
block_7_depthwise_... (BatchNormalizatio...	(None, 8, 8, 384)	0	block_7_depthwis...
block_7_depthwise_... (ReLU)	(None, 8, 8, 384)	24,576	block_7_depthwis...
block_7_project (Conv2D)	(None, 8, 8, 64)	256	block_7_project[...
block_7_project_BN (BatchNormalizatio...	(None, 8, 8, 64)	0	block_6_project_... block_7_project_...
block_7_add (Add)	(None, 8, 8, 64)		

block_8_expand (Conv2D)	(None, 8, 8, 384)	24,576	block_/_add[0][0]
block_8_expand_BN (BatchNormalizatio...	(None, 8, 8, 384)	1,536	block_8_expand[0...
block_8_expand_relu (ReLU)	(None, 8, 8, 384)	0	block_8_expand_B...
block_8_depthwise (DepthwiseConv2D)	(None, 8, 8, 384)	3,456	block_8_expand_r...
block_8_depthwise_... (BatchNormalizatio...	(None, 8, 8, 384)	1,536	block_8_depthwis...
block_8_depthwise_... (ReLU)	(None, 8, 8, 384)	0	block_8_depthwis...
block_8_project (Conv2D)	(None, 8, 8, 64)	24,576	block_8_depthwis...
block_8_project_BN (BatchNormalizatio...	(None, 8, 8, 64)	256	block_8_project[...
block_8_add (Add)	(None, 8, 8, 64)	0	block_7_add[0][0... block_8_project_...
block_9_expand (Conv2D)	(None, 8, 8, 384)	24,576	block_8_add[0][0]
block_9_expand_BN (BatchNormalizatio...	(None, 8, 8, 384)	1,536	block_9_expand[0...
block_9_expand_relu (ReLU)	(None, 8, 8, 384)	0	block_9_expand_B...
block_9_depthwise (DepthwiseConv2D)	(None, 8, 8, 384)	3,456	block_9_expand_r...
block_9_depthwise_... (BatchNormalizatio...	(None, 8, 8, 384)	1,536	block_9_depthwis...
block_9_depthwise_... (ReLU)	(None, 8, 8, 384)	0	block_9_depthwis...
block_9_project (Conv2D)	(None, 8, 8, 64)	24,576	block_9_depthwis...
block_9_project_BN (BatchNormalizatio...	(None, 8, 8, 64)	256	block_9_project[...
block_9_add (Add)	(None, 8, 8, 64)	0	block_8_add[0][0... block_9_project_...

block_10_expand (Conv2D)	(None, 8, 8, 384)	24,576	block_9_add[0][0]
block_10_expand_BN (BatchNormalizatio...	(None, 8, 8, 384)	1,536	block_10_expand[...
block_10_expand_re... (ReLU)	(None, 8, 8, 384)	0	block_10_expand_...
block_10_depthwise (DepthwiseConv2D)	(None, 8, 8, 384)	3,456	block_10_expand_...
block_10_depthwise... (BatchNormalizatio...	(None, 8, 8, 384)	1,536	block_10_depthwi...
block_10_depthwise... (ReLU)	(None, 8, 8, 384)	0	block_10_depthwi...
block_10_project (Conv2D)	(None, 8, 8, 96)	36,864	block_10_depthwi...
block_10_project_BN (BatchNormalizatio...	(None, 8, 8, 96)	384	block_10_project...
block_11_expand (Conv2D)	(None, 8, 8, 576)	55,296	block_10_project...
block_11_expand_BN (BatchNormalizatio...	(None, 8, 8, 576)	2,304	block_11_expand[...
block_11_expand_re... (ReLU)	(None, 8, 8, 576)	0	block_11_expand_...
block_11_depthwise (DepthwiseConv2D)	(None, 8, 8, 576)	5,184	block_11_expand_...
block_11_depthwise... (BatchNormalizatio...	(None, 8, 8, 576)	2,304	block_11_depthwi...
block_11_depthwise... (ReLU)	(None, 8, 8, 576)	0	block_11_depthwi...
block_11_project (Conv2D)	(None, 8, 8, 96)	55,296	block_11_depthwi...
block_11_project_BN (BatchNormalizatio...	(None, 8, 8, 96)	384	block_11_project...
block_11_add (Add)	(None, 8, 8, 96)	0	block_10_project... block_11_project...
block_12_expand (Conv2D)	(None, 8, 8, 576)	55,296	block_11_add[0][...

block_12_expand_BN (BatchNormalizatio...	(None, 8, 8, 576)	2,304	block_12_expand[...
block_12_expand_re... (ReLU)	(None, 8, 8, 576)	0	block_12_expand_...
block_12_depthwise (DepthwiseConv2D)	(None, 8, 8, 576)	5,184	block_12_expand_...
block_12_depthwise... (BatchNormalizatio...	(None, 8, 8, 576)	2,304	block_12_depthwi...
block_12_depthwise... (ReLU)	(None, 8, 8, 576)	0	block_12_depthwi...
block_12_project (Conv2D)	(None, 8, 8, 96)	55,296	block_12_depthwi...
block_12_project_BN (BatchNormalizatio...	(None, 8, 8, 96)	384	block_12_project...
block_12_add (Add)	(None, 8, 8, 96)	0	block_11_add[0][... block_12_project...
block_13_expand (Conv2D)	(None, 8, 8, 576)	55,296	block_12_add[0][...
block_13_expand_BN (BatchNormalizatio...	(None, 8, 8, 576)	2,304	block_13_expand[...
block_13_expand_re... (ReLU)	(None, 8, 8, 576)	0	block_13_expand_...
block_13_pad (ZeroPadding2D)	(None, 9, 9, 576)	0	block_13_expand_...
block_13_depthwise (DepthwiseConv2D)	(None, 4, 4, 576)	5,184	block_13_pad[0][...
block_13_depthwise... (BatchNormalizatio...	(None, 4, 4, 576)	2,304	block_13_depthwi...
block_13_depthwise... (ReLU)	(None, 4, 4, 576)	0	block_13_depthwi...
block_13_project (Conv2D)	(None, 4, 4, 160)	92,160	block_13_depthwi...
block_13_project_BN (BatchNormalizatio...	(None, 4, 4, 160)	640	block_13_project...
block_14_expand (Conv2D)	(None, 4, 4, 960)	153,600	block_13_project...

block_14_expand_BN (BatchNormalizatio...	(None, 4, 4, 960)	3,840	block_14_expand[...
block_14_expand_re... (ReLU)	(None, 4, 4, 960)	0	block_14_expand_...
block_14_depthwise (DepthwiseConv2D)	(None, 4, 4, 960)	8,640	block_14_expand_...
block_14_depthwise... (BatchNormalizatio...	(None, 4, 4, 960)	3,840	block_14_depthwi...
block_14_depthwise... (ReLU)	(None, 4, 4, 960)	0	block_14_depthwi...
block_14_project (Conv2D)	(None, 4, 4, 160)	153,600	block_14_depthwi...
block_14_project_BN (BatchNormalizatio...	(None, 4, 4, 160)	640	block_14_project...
block_14_add (Add)	(None, 4, 4, 160)	0	block_13_project... block_14_project...
block_15_expand (Conv2D)	(None, 4, 4, 960)	153,600	block_14_add[0][...
block_15_expand_BN (BatchNormalizatio...	(None, 4, 4, 960)	3,840	block_15_expand[...
block_15_expand_re... (ReLU)	(None, 4, 4, 960)	0	block_15_expand_...
block_15_depthwise (DepthwiseConv2D)	(None, 4, 4, 960)	8,640	block_15_expand_...
block_15_depthwise... (BatchNormalizatio...	(None, 4, 4, 960)	3,840	block_15_depthwi...
block_15_depthwise... (ReLU)	(None, 4, 4, 960)	0	block_15_depthwi...
block_15_project (Conv2D)	(None, 4, 4, 160)	153,600	block_15_depthwi...
block_15_project_BN (BatchNormalizatio...	(None, 4, 4, 160)	640	block_15_project...
block_15_add (Add)	(None, 4, 4, 160)	0	block_14_add[0][...] block_15_project...
block_16_expand (Conv2D)	(None, 4, 4, 960)	153,600	block_15_add[0][...

block_16_expand_BN (BatchNormalizatio...	(None, 4, 4, 960)	3,840	block_16_expand[...
block_16_expand_re... (ReLU)	(None, 4, 4, 960)	0	block_16_expand_...
block_16_depthwise (DepthwiseConv2D)	(None, 4, 4, 960)	8,640	block_16_expand_...
block_16_depthwise... (BatchNormalizatio...	(None, 4, 4, 960)	3,840	block_16_depthwi...
block_16_depthwise... (ReLU)	(None, 4, 4, 960)	0	block_16_depthwi...
block_16_project (Conv2D)	(None, 4, 4, 320)	307,200	block_16_depthwi...
block_16_project_BN (BatchNormalizatio...	(None, 4, 4, 320)	1,280	block_16_project...
Conv_1 (Conv2D)	(None, 4, 4, 1280)	409,600	block_16_project...
Conv_1_bn (BatchNormalizatio...	(None, 4, 4, 1280)	5,120	Conv_1[0][0]
out_relu (ReLU)	(None, 4, 4, 1280)	0	Conv_1_bn[0][0]
global_average_poo... (GlobalAveragePool...	(None, 1280)	0	out_relu[0][0]
dense_3 (Dense)	(None, 32)	40,992	global_average_p...

Total params: 2,298,976 (8.77 MB)

Trainable params: 40,992 (160.12 KB)

Non-trainable params: 2,257,984 (8.61 MB)

Training Transfer Learning model...

Epoch 1/20

42/42 ————— 20s 268ms/step - accuracy: 0.3789 - loss: 2.6536 -

Epoch 2/20

42/42 ————— 1s 19ms/step - accuracy: 0.9477 - loss: 0.4005 - va

Epoch 3/20

42/42 ————— 1s 19ms/step - accuracy: 0.9761 - loss: 0.1792 - va

Epoch 4/20

42/42 ————— 1s 20ms/step - accuracy: 0.9866 - loss: 0.1216 - va

Epoch 5/20

42/42 ————— 1s 19ms/step - accuracy: 0.9925 - loss: 0.0897 - va

Epoch 6/20

```

42/42 ————— 1s 19ms/step - accuracy: 0.9953 - loss: 0.0607 - va
Epoch 7/20
42/42 ————— 1s 19ms/step - accuracy: 0.9971 - loss: 0.0506 - va
Epoch 8/20
42/42 ————— 1s 24ms/step - accuracy: 0.9993 - loss: 0.0392 - va
Epoch 9/20
42/42 ————— 1s 22ms/step - accuracy: 1.0000 - loss: 0.0320 - va
Epoch 10/20
42/42 ————— 1s 24ms/step - accuracy: 1.0000 - loss: 0.0265 - va
Epoch 11/20
42/42 ————— 1s 18ms/step - accuracy: 1.0000 - loss: 0.0247 - va
Epoch 12/20
42/42 ————— 1s 19ms/step - accuracy: 1.0000 - loss: 0.0199 - va
Epoch 13/20
42/42 ————— 1s 18ms/step - accuracy: 1.0000 - loss: 0.0181 - va
Epoch 14/20
42/42 ————— 1s 20ms/step - accuracy: 1.0000 - loss: 0.0157 - va
Epoch 15/20
42/42 ————— 1s 19ms/step - accuracy: 1.0000 - loss: 0.0139 - va
Epoch 16/20
42/42 ————— 1s 19ms/step - accuracy: 1.0000 - loss: 0.0124 - va
Epoch 17/20
42/42 ————— 1s 18ms/step - accuracy: 1.0000 - loss: 0.0115 - va
Epoch 18/20
42/42 ————— 1s 20ms/step - accuracy: 1.0000 - loss: 0.0113 - va
Epoch 19/20
42/42 ————— 1s 20ms/step - accuracy: 1.0000 - loss: 0.0096 - va
Epoch 20/20
42/42 ————— 1s 20ms/step - accuracy: 1.0000 - loss: 0.0099 - va

```

Evaluating Transfer Learning model on test data...

Transfer Learning Test Loss: 0.0568

Transfer Learning Test Accuracy: 0.9930

--- Accuracy Comparison (CNN vs LSTM vs Transfer Learning) ---

CNN Test Accuracy: 0.8780

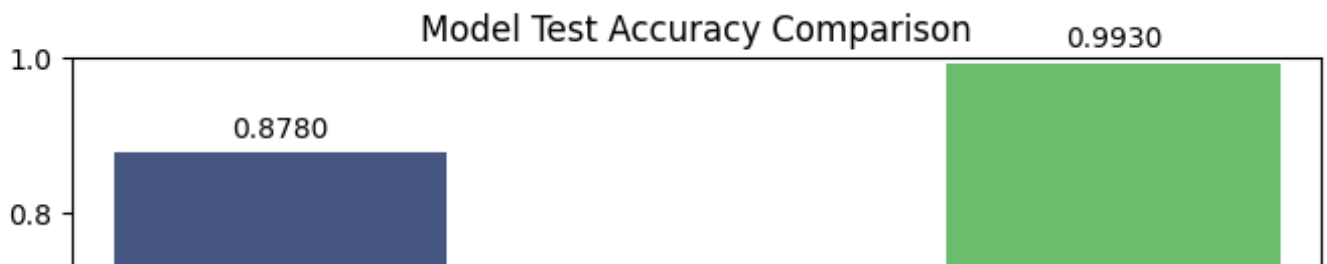
LSTM Test Accuracy: 0.0383

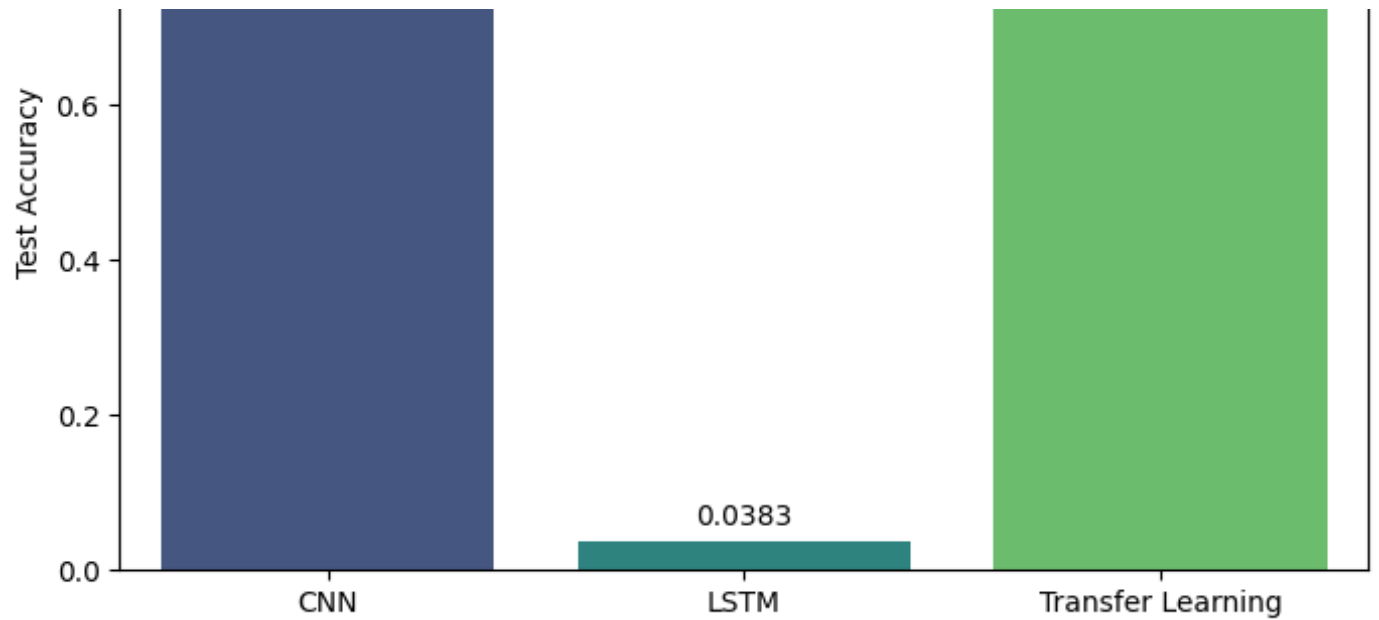
Transfer Learning Test Accuracy: 0.9930

/tmp/ipython-input-17-585755722.py:97: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in

```
sns.barplot(x=labels_comp, y=accuracies_comp, palette='viridis')
```





--- Loss Comparison (CNN vs LSTM vs Transfer Learning) ---

CNN Test Loss: 0.6464

LSTM Test Loss: 3.4529

Transfer Learning Test Loss: 0.0568

/tmp/ipython-input-17-585755722.py:114: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in

```
sns.barplot(x=labels_comp, y=losses_comp, palette='magma')
```







--- Transfer Learning Evaluation ---
Generating Classification Report and Confusion Matrix for Transfer Learning Mc
9/9 7s 443ms/step

Transfer Learning Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	9
1	0.91	1.00	0.95	10
2	1.00	1.00	1.00	10
3	1.00	1.00	1.00	11
4	1.00	1.00	1.00	11
5	1.00	1.00	1.00	8
6	1.00	1.00	1.00	10
7	1.00	1.00	1.00	8
8	1.00	1.00	1.00	8
9	1.00	1.00	1.00	9
10	1.00	1.00	1.00	7
11	1.00	1.00	1.00	10
12	1.00	1.00	1.00	8
13	1.00	1.00	1.00	10
14	1.00	1.00	1.00	9
15	1.00	0.88	0.93	8
16	1.00	1.00	1.00	12
17	1.00	1.00	1.00	10
18	1.00	1.00	1.00	9
19	1.00	1.00	1.00	10
20	1.00	1.00	1.00	9
21	1.00	1.00	1.00	8
22	1.00	1.00	1.00	8
23	0.91	1.00	0.95	10
24	1.00	0.88	0.93	8
25	1.00	1.00	1.00	8
26	1.00	1.00	1.00	8
27	1.00	1.00	1.00	8
28	1.00	1.00	1.00	8
29	1.00	1.00	1.00	9
30	1.00	1.00	1.00	8
31	1.00	1.00	1.00	8
accuracy			0.99	287
macro avg	0.99	0.99	0.99	287
weighted avg	0.99	0.99	0.99	287





#

✓ Predict with Transfer learning model(MobileNetV2)

#

```

import matplotlib.pyplot as plt
import numpy as np
from tensorflow.keras.models import Model
from tensorflow.keras.layers import GlobalAveragePooling2D
from tensorflow.keras.applications import MobileNetV2 # Import MobileNetV2

# Check if the transfer_model and test data are available
if 'transfer_model' in locals() and 'X_test_rgb' in locals() and 'y_test' in locals():
    print("\n--- Predicting and Plotting Samples with Transfer Learning Model ---")

    num_samples_to_plot = 5 # Number of random samples to plot
    sample_indices = random.sample(range(X_test_rgb.shape[0]), min(num_samples_to_plot, X_test_rgb.shape[0]))

    plt.figure(figsize=(15, 5 * ((num_samples_to_plot + 4) // 5))) # Adjust figure size

    for i, sample_index in enumerate(sample_indices):
        # Get the sample image and true label
        sample_image = X_test_rgb[sample_index]
        true_label_encoded = y_test[sample_index]
        true_label_index = np.argmax(true_label_encoded)

        # Predict the class for the sample image
        sample_image_input = np.expand_dims(sample_image, axis=0) # Add batch dimension
        prediction_probs = transfer_model.predict(sample_image_input)
        predicted_label_index = np.argmax(prediction_probs)

```

```

confidence = np.max(prediction_probs) * 100

# Get the original species label from the unique_labels list
# Ensure unique_labels is available
if 'unique_labels' in locals():
    true_species = unique_labels[true_label_index]
    predicted_species = unique_labels[predicted_label_index]
else:
    true_species = f"Class {true_label_index}"
    predicted_species = f"Class {predicted_label_index}"
    print("Warning: unique_labels not found, using integer indices for species")

display_image = (sample_image[:, :, 0] * 255).astype(np.uint8)

plt.subplot(1, num_samples_to_plot, i + 1) # Arrange plots in a row
plt.imshow(display_image, cmap='gray')
plt.title(f"True: {true_species}\nPred: {predicted_species} (confidence: {confidence})")
plt.axis('off')

plt.tight_layout()
plt.show()

else:
    print("Transfer learning model or test data not available. Cannot predict and")
    print(f"Is transfer_model defined: {'transfer_model' in locals()}")
    print(f"Is X_test_rgb defined: {'X_test_rgb' in locals()}")
    print(f"Is y_test defined: {'y_test' in locals()}")
    if 'X_test_rgb' in locals():
        print(f"X_test_rgb shape: {X_test_rgb.shape}")
    if 'y_test' in locals():
        print(f"y_test shape: {y_test.shape}")

```



--- Predicting and Plotting Samples with Transfer Learning Model ---

1/1  2s 2s/step
 1/1  0s 34ms/step
 1/1  0s 35ms/step
 1/1  0s 36ms/step
 1/1  0s 36ms/step

True: 23
Pred: 23 (99.9%)



True: 6
Pred: 6 (97.2%)



True: 30
Pred: 30 (100.0%)



True: 31
Pred: 31 (100.0%)



True: 17
Pred: 17 (99.9%)



Conclusion: Implementing Transfer Learning for Leaf Classification

Transfer learning was applied using **MobileNetV2** as the base model to improve classification performance on the Flavia leaf dataset.

Input Data Configuration for Transfer Learning

Dataset	Shape
Training Set	(1334, 128, 128, 3)
Validation Set	(286, 128, 128, 3)
Test Set	(287, 128, 128, 3)

- All images were converted to **RGB** and resized to **128×128×3**.
- **MobileNetV2 weights** were loaded from TensorFlow's Keras application hub (excluding top layers).

Transfer Learning Model Summary

- Base model: **MobileNetV2** (frozen layers)
- Top layers added:
 - GlobalAveragePooling2D
 - Dense(32, activation='softmax')

Parameters	Count
Total Parameters	2,298,976
Trainable Parameters	40,992
Non-Trainable Parameters	2,257,984

Training Results

The model was trained for **20 epochs** with **early stopping** based on validation loss.

Key Training Metrics (Final Epoch):

- **Training Accuracy:** 100%
- **Validation Accuracy:** 98.6%
- **Validation Loss:** 0.0514

Final Test Performance

Metric	Value
Test Accuracy	99.3%
Test Loss	0.0568

Model Comparison Summary

Model	Test Accuracy	Test Loss
CNN (custom)	87.8%	0.6464
LSTM	3.8%	3.4529
Transfer Learning (MobileNetV2)	99.3%	0.0568

Transfer learning outperformed both CNN and LSTM by a significant margin.

Transfer Learning Classification Report

- **Overall Accuracy:** 99.3%
- **Macro Average F1-Score:** 0.99
- **Weighted Average F1-Score:** 0.99
- Most classes achieved perfect precision, recall, and F1-score (1.00).
- A few classes had slight misclassifications with F1-scores around 0.93–0.95.

Notable Misclassifications:

- Class 15: F1-score = 0.93
- Class 24: F1-score = 0.93
- Class 1 and Class 23: F1-score = 0.95

These may indicate subtle leaf variations or overlaps in morphological features.

Key Takeaways

- **Transfer Learning with MobileNetV2** dramatically improved performance while requiring fewer trainable parameters.
 - Even with limited training data (~1,300 images), leveraging a pretrained network proved highly effective.
 - The **LSTM model** was not suitable for spatial image data and severely underperformed.
 - **CNN performed decently**, but fell short compared to the transfer learning approach.
-

END

END

